

Deliverable report

“Alessio Blascovich”: Mat: 257561, alessio.blascovich@studenti.unitn.it, *GitRepo*:
<https://github.com/elblasco/GPU-Computing-2025-257561>

Abstract—The sparse matrix-dense vector multiplication (SpMV) is an algebraic operation in which a sparse matrix is multiplied by a dense vector. SpMV is widely used in real-world applications such as scientific and engineering computing, data mining, and graph analysis.

This deliverable discusses the implementation and the performance analysis of a series of SpMV algorithms for COO matrices. These algorithms are implemented for a single-core CPU and a CUDA capable GPU. Regarding the GPU implementations, no threads/warp level synchronisation has been used, only atomic operations to keep the resulting vector consistent.

Index Terms—Sparse Matrix, SpMV, CUDA, Parallelization, Storage Format

I. INTRODUCTION

SpMV is considered a fundamental algebraic operation in many scientific fields such as computational fluid dynamics, optimisation and graph algorithms.

The main challenges in implementing an efficient SpMV algorithm differ in CPU and GPU, and are strictly bound to the type of matrix representation used (COO, CSR, ELL,...). For the purpose of this deliverable only COO representation is taken into account.

From the CPU point of view, the main issue is to maximise the cache hits throughout the execution, data consistency is not a problem since only single core CPUs were used.

On the GPU side the problem of maximising cache hits remains, and on top of that, race conditions are added. Since in COO format the length of a row is not specified, multiple threads could be reading different portions of the same row and therefore, writing on the same output memory area.

II. PROBLEM STATEMENT

The SpMV operation is formally defined in Equation 1, where A is a $n \times m$ sparse matrix, x a $m \times 1$ dense vector and y another $n \times 1$ dense vector.

$$y = Ax \quad (1)$$

A. Storage Format

The coordinate (COO) format is a representation of sparse matrix. COO uses 3 different arrays to store all the necessary information. One array dedicated to store all the non-zero values, while the other 2 contain the corresponding row and column indices. For the sake of simplicity, all the input matrices are assumed to be sorted in row-major order, this is done by transposing all the input matrices which were sorted in column-major order.

B. Parallelization

2 different kernels have been used for the CUDA capable GPU, although in both of them the workload is divided equally among all the threads used. The main difference is the pattern used to access the global memory from a thread. In the first approach each thread jumps of stride equivalent to the number of threads used, while in the second approach a thread access to a continuous area of memory long as the total number of non-zero element in the matrix divided by the number of threads used.

III. STATE OF THE ART

Starting from the implementations presented in this deliverable, performances could be drastically enhanced. Fukaya et al. [1], despite focussing on the HDC storage format, improved the SpMV CPU based by exploiting diagonalization patterns, SIMD instructions and *OpenMP* directives. On the GPU side Chu et al. [2] used the CSR format and a more sophisticated load balancing to improve the performances. NVIDIA itself made available a set of libraries to manage and use COO formatted matrices, namely *cuSPARSE*.

IV. METHODOLOGY AND CONTRIBUTIONS

The tests has been carried out performing multiple rounds of caches warm up and then executing the various algorithm after. The data have been collected throughout the execution since only for the first warm up cycle there were substantially data variations.

A. CPU implementation

The CPU simply loop through all the elements and performs the operations required for the matrix multiplication element-wise. The CPU Algorithm 1 maximise the cache hit since it reads, sequentially, elements that are stored in a contiguous memory area, this holds for `res[row[i]]` and `val[i]`.

Algorithm 1 SpMV algorithm on CPU

Require: The input vectors *val*, *row* and *col* of size *nnz*, array *arr* of size *columns* and an array *res* of size *rows*.

```
1: procedure FUNCTION(row, col, val, arr, res, ...)
2:   while i in {0...nnz} do
3:     res[row[i]] = res[row[i]] + (val[i] · arr[col[i]])
4:     i = i + 1
5:   end while
6: end procedure
```

In particular instance I had to use my personal computer since the cluster version of *valgrind* has some trouble with

the cache profile tool. Due to hardware limitation I performed the cache profiling using smaller matrices. To have clearer I did not perform any warm-up cycle sine I want to measure the effectiveness of my algorithm on “pseudo-empty” caches.

Matrix	D1 read	D1 write
lp_nug05	1.5%	.5%
rdb2501	.8%	.3%
fpga_dcop_15	.3%	.2%

TABLE I
L1 CACHE MISSES

Matrix	LLd read	LLd write
lp_nug05	.7%	.3%
rdb2501	.4%	.2%
fpga_dcop_15	.1%	.1%

TABLE II
L2 AND L3 CACHE MISSES

B. GPU

For the GPU implementation I have created two kernel, they differ on the access pattern used for the memory. In Algorithm 2 every thread jumps to the next element which has index $current\ index + numThreads$. In Algorithm 3 each thread accesses $portion$ elements which are stored in a contiguous area of memory.

Algorithm 2 SpMV algorithm with striding

Require: The input vectors val , row and col of size nnz , array arr of size $columns$, an array res of size $rows$ and a number $portion$.

```

1: procedure FUNCTION( $row, col, val, arr, res, \dots$ )
2:    $start \leftarrow threadId$ 
3:    $end \leftarrow (portion \cdot numThreads) + start$ 
4:   while  $i$  in  $\{start \dots \min\{end, nnz\}\}$  do
5:      $atomicAdd(\&res[row[i]], val[i] \cdot arr[col[i]])$ 
6:      $i = i + numThreads$ 
7:   end while
8: end procedure
```

Algorithm 3 SpMV algorithm with contiguous chunk

Require: The input vectors val , row and col of size nnz , array arr of size $columns$, an array res of size $rows$ and a number $portion$.

```

1: procedure FUNCTION( $row, col, val, arr, res, \dots$ )
2:    $start \leftarrow portio \cdot threadId$ 
3:    $end \leftarrow portion + start$ 
4:   while  $i$  in  $\{start \dots \min\{end, nnz\}\}$  do
5:      $atomicAdd(\&res[row[i]], val[i] \cdot arr[col[i]])$ 
6:      $i = i + 1$ 
7:   end while
8: end procedure
```

From the theory Algorithm 2 should have been faster. This follows from the fact that every thread within the same SM (Streaming Multiprocessor) accesses elements that are spacially adjacent in the memory. Contrary, Algorithm 3 has an higher rate of cache miss since, threads in the same SM may need to access elements far away from each other.

V. SYSTEM DESCRIPTION AND EXPERIMENTAL SET-UP

A. System Description

Table III shows the various software used in the deliverable, excluded the building tools. It is worth mentioning again that, for *valgrind*, I had to do a rollback on my pc since the last version (same of the cluster) has some problem with the tool *cachegrind*.

System	GCC	Valgrind	CUDA
Cluster	11.5.0	3.23.0	12.5.0
PC	15.1.1	3.20.0	12.8.1

TABLE III
SYSTEMS SOFTWARES

System	Processor	Cores per Socket	RAM	Accelerator
Cluster	Intel(R) Xeon(R) Silver 4214	48 at 2.2 GHz	400 GB	NVidia A30
PC	AMD Ryzen 5 3500U	8 at 2.1 GHz	8 GB	Radeon Vega Mobile Gfx

TABLE IV
SYSTEMS DETAILS

B. Dataset description

To add a bit more of generality to the dataset I picked four matrix from four different kinds from the site <https://sparse.tamu.edu/>. Since all the matrices are squared only the number is written on Table V.

- 1) Model Reduction Problem;
- 2) Electromagnetics Problem;
- 3) Undirected Weighted Graph;
- 4) Optimization Problem.

Name	Rows	Pattern Entries	Symmetric	Kind
bone010	986,703	71,666,325	Yes	1
dielFilterV3real	1,102,824	89,306,020	Yes	2
mawi_201512020330	226,196,185	480,047,894	Yes	3
nlpkkt160	8,345,600	229,518,112	Yes	4

TABLE V
INPUT MATRICES

C. Experimental Set-up

I used 5 warm-up runs followed by 5 runs to collect data. I kept this set-up even though after few tests one warm-up round was enough to have stable values. Each job submitted to the cluster executed a full warm-up and run for Algorithm 1, then Algorithm 2 and then Algorithm 3.

Average execution time




Fig. 1. Average Execution Time (lower is better)

Average GFLOP/s developed




Fig. 2. Average GFLOP/s Developed (lower is better)

VI. EXPERIMENTAL RESULTS

As expected in Section IV and shown in Figure [1, 2], Algorithm 2 is faster than all the other implementations.

I decided to dig a bit more into the performances of Algorithm 2 by changing the maximum amount of thread per block. Starting from the base of 512 threads I divided by 2 until reaching 32, i.e., the size of warp. The optimal amount of threads per block resulted to be 256. This is due to the use of an atomic instruction within the code, the semantic of the instruction make concurrent write over the same address sequential; in this case if two threads are working on the same row their sum operations will be sequentialised since they are trying to write on the same result address. Using less threads, i.e., each thread works on more elements; conflicts (threads working) on the same row are reduced to a minimum. However, this minimum is not stable, in fact decreasing again the value to 128, 64 or 32 the performances decrease again, this could be a symptom of either:

- too few threads working, so an excessive workload for each of them;
- too much threads working on the same row, hence a high number of conflicts on the same row.

Unfortunately, I was not able to determine a more fine grain optimisation, since the optimal size depends on the input ma-

GFLOP/s changing threads per block




Fig. 3. GFLOP Compared The Number Of Threads Per Block

GB/s changing threads per block




Fig. 4. GB/s Compared The Number Of Threads Per Block

trix shape. Figure [3, 4] highlights the fact that the variations of performances are not constant over all the matrices considered.

VII. CONCLUSIONS

[max 200 words]

Summarize findings and future work ...

REFERENCES

- [1] T. Fukaya, K. Ishida, A. Miura, T. Iwashita, and H. Nakashima, "Accelerating the spmv kernel on standard cpus by exploiting the partially diagonal structures," 2021. [Online]. Available: <https://arxiv.org/abs/2105.04937>
- [2] G. Chu, Y. He, L. Dong, Z. Ding, D. Chen, H. Bai, X. Wang, and C. Hu, "Efficient algorithm design of optimizing spmv on gpu," in *Proceedings of the 32nd International Symposium on High-Performance Parallel and Distributed Computing*, ser. HPDC '23. New York, NY, USA: Association for Computing Machinery, 2023, p. 115–128. [Online]. Available: <https://doi.org/10.1145/3588195.3593002>